

How we measured device and AI policies in UT syllabi

Status: complete. All 4,089 classifiable documents are coded and loaded. Two cleanup items remain open and are flagged where they affect a number; see [docs/known-issues.md](#).

This document explains where the data came from, how each syllabus was scored, what we did to check the scores, and what the data cannot support. It is written to be argued with. If you think a number is wrong, the last section tells you exactly how to prove it.

1. The short version

We collected every syllabus UT published for Fall 2026 — 4,094 documents covering 6,968 course sections and 2,475 instructors — converted 4,089 of them to searchable text (§4 accounts for the five), and scored each one on three separate 0-5 scales: how tightly it controls **phones**, how tightly it controls **laptops**, and how tightly it controls **generative AI**.

The scoring was done by a language model working from a written codebook, in batches of 25 documents, with no knowledge of the results of any other batch.

The scoring is not asked to be trusted. Every score above zero carries one or more quotes pulled verbatim from the syllabus, and every quote was checked by a program, character by character, against the stored text of that syllabus before the score was allowed into the database. A score whose quote does not literally appear in the document is rejected and never loaded. So the underlying claim of this dataset is not “a model judged this syllabus strict.” It is: *here is the sentence the instructor wrote; here is the bin we put that sentence in; go read it.*

That is a stronger evidentiary standard than most hand-coding projects meet, because human coders rarely record the sentence that drove their decision.

2. Where the documents came from

Texas Education Code §51.974 requires public universities to post syllabi for undergraduate courses on a publicly accessible website, no login required. UT does this at:

<https://utdirect.utexas.edu/apps/student/coursedocs/courses/nlogin/>

We queried all 282 departments in the site’s own dropdown for Fall 2026, in residence. Everything we collected was already public and free to read; we added no credentials and bypassed nothing. The technical quirks of the site are documented in `docs/source-notes.md`.

Syllabi arrive through two channels, and both were collected:

Channel	How it works	Documents
Direct download	A PDF or Word file at <code>/download/<id>/</code>	3,282
Simple Syllabus	A web app that renders the syllabus in the browser	812

The second channel is the one worth flagging: those syllabi are not files. They are pages built by JavaScript after the page loads, so an ordinary download gets an empty shell. We rendered each one in a real headless browser and saved the resulting text. Had we skipped them, we would have silently dropped a fifth of the corpus, and not at random — Simple Syllabus adoption varies by college, so the missing fifth would have been concentrated in particular schools.

Deduplication. Files are stored under the SHA-256 hash of their contents, so a syllabus posted under several course IDs is stored once and linked to each. This is why the document count (4,094) is smaller than the section count (6,968): one document often serves many sections of the same course.

3. The unit-of-analysis problem

This matters more than any other methodological choice here, because it is the one most likely to produce a wrong number in print.

A syllabus is not a course, and a course is not an instructor. The same PDF may cover eleven sections. One instructor may post four different syllabi. So “38% of syllabi restrict phones” and “38% of classes restrict phones” are different claims, and only the first is what we measured.

Concretely: 1,383 **documents** mention ChatGPT, but those documents cover 2,259 **sections**. Reporting documents as if they were classes would be off by roughly 900.

Three denominators are available, and the story should name which one it uses:

Denominator	Corpus	Coded	What a percentage of it means
Documents	4,094	4,089	share of posted syllabi
Sections	6,968	6,963	share of scheduled class sections
Instructors	2,475	2,473	share of instructors

The coded column is the denominator used in §9. It is smaller because four encrypted files yielded no text and one document was quarantined (§4).

An early measurement of one indicator moved from 38.2% to 39.7% to 43.2% across these three — same underlying data, three defensible sentences. **Recommended default: sections**, because “what share of UT classes ban phones” is the question readers actually have, and because it does not over-weight departments that post one templated syllabus across many sections.

Near-duplicates, measured. Content hashing catches only byte-identical files. Syllabi differing by a section number, a room or a date hash differently and survive as separate documents. `scripts/find_near_duplicates.py` groups them by word-shingle similarity: **354 clusters covering 978 documents, of which 624 are redundant copies — 15.3% of the corpus.** The largest are a 27-copy rhetoric template, a 22-copy Spanish template and a 13-copy cello studio family.

Collapsing them moves the headline figures by **−0.7 to +1.3 points**, so corpus-level percentages stand as reported. Department-level breakouts do not: templating is concentrated in particular departments, so a Spanish or rhetoric breakout computed per document would be badly over-weighted. Compute those per cluster, using the `near_duplicate_group` column.

4. Getting the text out

Nothing can be scored that was not read correctly, so this stage got its own verification.

Method	Documents	Notes
Existing text layer	3,257	Read directly from the PDF
Rendered from web app	812	Simple Syllabus, captured from the browser
OCR	20	Scanned or image-only PDFs
Encrypted, not recovered	4	Password-protected; decryption failed, no text
Quarantined	1	See below
Total	4,094	of which 4,089 yielded text and were coded

The four encrypted files are the only genuine extraction losses: 0.1% of the corpus. They are counted here rather than quietly dropped, and they are excluded from every percentage in §9 because there is nothing in them to score.

Two failure modes were found and fixed during this stage, and both are the kind that would have quietly corrupted results:

Text that isn't text. Nine PDFs had a text layer that passed every size-based sanity check but decoded to gibberish — the font was subset with a scrambled character map, so the file “contained text” that was really a cipher. A character-count check cannot catch this. We added a check that looks for common short words, and re-OCR'd the failures. Eight recovered cleanly. The ninth was a French-language syllabus that our English-only word list had falsely accused, which is why the check now covers six languages.

Columns read across, not down. 71 PDFs are laid out in multiple columns, and the default extraction setting interleaves the columns line by line, producing text that looks fine but reads as nonsense across column boundaries. These are recorded in `data/multicolumn_ids.json` and re-extracted with column-aware settings. (*Cleanup pass in progress.*)

One document was removed from the corpus. Document 2040 is not a syllabus. It is a 14 MB file of personnel records — roughly 43 job interviews for a College of Liberal Arts career-coach position, with named candidates and written assessments of each. It was published on UT's public syllabus site, where it remains live. We removed its text from our database, deleted our copy of the text, and logged the removal in a `quarantined_documents` table rather than deleting the record silently. A scan of the full corpus found no other document of this kind. **An open editorial question: whether to notify UT that the file is exposed, and whether the exposure is itself a story.**

5. What was actually measured

Three scales, scored independently, 0-5. Zero always means the syllabus does not address that thing; 5 is the tightest control. The full anchors are in `prompts/policy-classification-v7.md`.

Alongside each scale, two fields that deliberately sit **off** the 0-5 axis:

- `<scale>_required` — whether students are told to *have* the thing. A syllabus can require a laptop and restrict its use in class; that is one coherent policy, and a single strictness scale can only hold half of it. Before this field existed, coders were discarding one of the two facts on roughly one document in twenty.
- `<scale>_status` — `clear`, `undecided`, or `contradictory`. This keeps apart two situations that look identical in a single number: an instructor who wrote two rules that genuinely conflict, and an instructor who posted an unedited template and never chose. As one coder wrote about a template case, the score has “no principled basis.” Those documents belong outside a strict-versus-lenient ratio, not averaged into it.

How the codebook was built

The instrument went through seven versions before the production run. Every revision is recorded in `prompts/README.md`, including what broke and why. Two are worth repeating here, because they are cases where a plausible-looking rule would have produced a badly wrong published number:

- **v2’s rule R1** said that any stated exception dropped an absolute ban from 5 to 4. But accommodation language is nearly universal: 95% of documents in this corpus contain the word “accommodation” and 75% name UT’s disability services office. Read literally, that rule would have emptied the strictest category almost completely — the exact bin the story is about. It was caught because a coder flagged the ambiguity instead of quietly picking a reading.
- **v2’s rule R2** banned inference so broadly that “Computer: PC or Mac” in a required-materials list scored as *silence on laptops*, discarding a syllabus that plainly does tell students to have a computer.

Neither error would have been visible in the output. Both were found by reading what the coders said they struggled with — which is why every batch is asked to report its hard calls, and why those reports are read.

The scoring was version-controlled

`document_policies.prompt_version` records which codebook version produced each row. Revised instructions add a comparable pass rather than overwriting the old one, so the size of a revision’s effect can be stated rather than asserted.

900 documents carry scores under both v6 and v7. Comparing them:

scale	scores changed	moved 2+ bins	strict share, v6 → v7
Phone	4.6%	20	75.8% → 79.0% (+3.2)
Laptop	4.2%	21	53.7% → 56.2% (+2.5)
AI	11.3%	12	36.7% → 46.4% (+9.8)

This is worth sitting with. The device scales barely moved — v7’s device changes were clarifications, and they behaved like clarifications. **The AI scale moved almost ten points** on the same 900 documents, because v7 rewrote the rule governing the university’s “partial basis” template, which a quarter of the corpus uses.

Two things follow. First, the instrument is sensitive to how one rule is worded on the scale where the language is newest and least settled, which is exactly what §7.7 says about the remaining band. Second, this is why a published figure has to name its codebook version: “half of AI policies are restrictive” and “a third of AI policies are restrictive” are both defensible sentences about the same syllabi under two versions of the same instrument, six weeks apart. The version is part of the claim, not a footnote to it.

6. How the scores are checked

Four independent checks. They test different things, and none of them requires taking the classifier's word for anything.

6.1 Every quote is machine-verified (audited: 11,497 quotes, 0 failures)

This is the backbone. Each score above zero must carry verbatim evidence, and a loader program checks each quote against the syllabus text as stored in the database. Quotes that do not match are **rejected and never loaded** — the record is skipped whole and the reason written to `data/rejected.jsonl`.

The check normalizes only what PDF extraction demonstrably breaks: smart quotes, zero-width characters, the non-breaking spaces UT boilerplate is full of, and line-break hyphenation. It does not normalize meaning and does not accept paraphrase. Coders are forbidden from stitching sentences together or reconstructing across column breaks. A looser second pass would accept a 0.92 similarity match to tolerate broken ligatures.

The audit. `scripts/verify_evidence.py` re-checks every stored quote against every stored document after the fact, without trusting the loader:

rows audited	4,089
quotes checked	11,497
matched exactly	11,497 (100.00%)
needed the looser pass	0
not found	0

Every quote matches its syllabus character for character. The looser pass exists but was never once required. Run the script to reproduce this.

This check caught a real error of ours. We re-extracted document 2352's text and updated the database but forgot to regenerate the copy the coder was reading. The coder scored 86 KB of mojibake that no longer existed anywhere. The validator caught it immediately, because it checks against the database rather than against the coder's input. `scripts/export_for_coding.py` now regenerates coder files from the database and reports any that were stale.

6.2 Inter-coder agreement (175 documents, measured)

Seven batches, selected by a recorded random seed before the run began, were coded a second time by an independent coder that never saw the first result.

Results. Krippendorff's alpha uses the ordinal difference function; exact and within-one agreement are also given, because alpha alone is hard to interpret.

scale	exact	within one bin	alpha	documents differing
Phone	96.6%	99.4%	0.982	6
Laptop	98.3%	99.4%	0.992	3
AI	98.3%	100.0%	0.995	3

The off-scale fields agree at 98.9-100%. Only **two** documents in 175 differ by more than one bin on any scale.

Independence was verified, not assumed. Two passes could agree because one saw the other. They did not: on 74% of documents the two coders chose *different verbatim quotes* to justify the same score, and 99% of their written notes differ entirely. The scores converge while the reasoning artifacts diverge, which is what genuine independent coding looks like.

What this number does and does not mean — read this before quoting it.

Alpha of 0.98 is higher than published content analysis usually reports, and that should raise an eyebrow rather than settle the question. The reason is that both coders are the same kind of reader applying the same instructions. Two human undergraduates bring different priors, different attention, and different tolerance for ambiguity; two runs of the same model bring far less of that variation. So this figure measures the **reproducibility** of the instrument - does the codebook return the same answer when applied again - and not its **validity**, which is whether that answer is correct.

A codebook can be perfectly reproducible and consistently wrong. If a rule mis-reads a common phrasing, both passes mis-read it identically and agreement stays at 0.98. That is precisely the correlated-error risk described in §8, and this statistic is blind to it by construction.

So the honest claim is narrow: the scores are not arbitrary, and re-running the process would produce materially the same dataset. Anyone who wants to know whether the scores are *right* has to look at §6.3, which is the only check that compares against a different kind of reader.

The disagreements are informative. Both wide ones sit exactly where coders reported the codebook is thin: document 251 (laptop 2 vs 0) turns on whether a passing mention of a device is a permission or silence, and document 1005 (phone 3 vs 5) on whether a politely worded directive with a soft hedge is a rule or a preference. Both are on the v8 list in `prompts/README.md`.

6.3 Sealed human validation set (200 documents)

200 documents were drawn at random and sealed **before any instructions were written**, and have never been used to develop or tune the codebook. Desk staff can code these by hand from the same written instructions and compare against the machine scores.

This is the check that answers “should we believe any of this,” and the only one that requires no trust in the classifier at all. It is deliberately sized so that a few people can finish it in an afternoon. **We recommend doing it before publication, and reporting the result whatever it is.**

6.4 Spot-checking by anyone, at any time

Because every score carries its quote and its source URL, any individual row can be verified in about thirty seconds by opening the syllabus and reading the sentence. There is no step where a reader has to accept a judgment they cannot inspect. An editor who wants to check twenty rows before signing off can do that.

7. What this data cannot tell you

Stated plainly, because these will be asked. Open defects and pending cleanup are tracked separately in `docs/known-issues.md`, with measured sizes.

1. **A syllabus is not a classroom.** We measured what instructors *wrote*, not what they enforce or what students experience. A strict written policy that is never enforced scores 5. This is a real and unfixable limitation of syllabus data, and the story should say so early rather than defend it later.
2. **We cannot measure change over time yet.** This is one semester. Any claim that policies got stricter “this year” needs a prior-year corpus, which we do not have. The site does serve past semesters, so this is collectable — but it is not collected, and the current data cannot support a trend sentence.
3. **Generic device wording blurs the phone/laptop distinction.** Many syllabi say “electronic devices” without naming phones or laptops. Rule R7 applies such a rule to both scales, which is the right reading of the text, but it means that when a document’s phone and laptop scores are identical, that may reflect *absence of specificity* rather than a deliberate matching policy. Do not report phone/laptop agreement rates without accounting for this. A three-bucket framing — names phones specifically / names laptops specifically / generic device language — is more honest than a two-way split.
4. **phone_required is inflated** by the same generic wording, and needs a correction factor before it is quoted.
5. **Exam-only rules are not class rules — and the codebook is not yet consistent about them.** A ban that applies only during exams leaves ordinary class meetings unaddressed, and scores 0 with a note. Scoring it otherwise would assert a restriction the instructor did not write — but scoring 0 without the note, as one coder put it, “asserts a silence that isn’t real.”

A sentence that is *both* a requirement and exam-scoped (“on exam days, bring a charged device with LockDown Browser installed”) hits two rules that point opposite ways, and coders split on it: of 78 affected documents coded so far, 58 scored the laptop requirement and 10 scored silence. This is the clearest example of the correlated-error risk described in §8, and it was found the way

§8 says such errors get found — seven coders raised it independently. It is tracked as issue 1 in `docs/known-issues.md`. **A laptop-requirement figure should not be published until it is resolved.**

6. **Smart glasses and display eyewear sit outside all three scales.** 200 documents (4.9%, covering 382 sections) mention them, and a new McCombs “Smart Glasses Course Policy” bans them outright. Because eyewear is neither a phone nor a laptop, those bans are recorded in `coder_notes` and score 0. **56 documents — about a third of the coded smart-glasses set — score 0 on both device scales while actually carrying a device ban**, so they read as silent when they are not. That is 1.4% of the corpus, and it understates the strict end rather than overstating it. They are scored 0 with a note, because they are neither phone nor laptop. This is arguably its own small story rather than a flaw.
7. **The AI conditional/restricted boundary carries about a 3-point band.** A quarter of the corpus uses UT’s “generative AI is permitted on a partial basis” template. The rule that sorts those documents asks whether the instructor named which tasks may use AI or deferred to later instruction, and it does discriminate — but it does not say which wins when a syllabus does both. 257 documents (7.6% of those stating an AI policy) are affected. Flipping all of them moves the strictness figure between **49.1% and 56.7%** against 50.9% as coded. Report the band; do not report a bare 51%. Tracked as issue 1 in `docs/known-issues.md`.
8. **The codebook is a set of choices, not a law of nature.** Where to draw the line between 3 and 4 is a judgment. It is written down, applied consistently, and versioned — but a different newsroom drawing the line differently would get different percentages from the same syllabi. The bins are in the repo; anyone is free to disagree with them in public.

8. On the use of a language model

Most of this desk is skeptical of AI-produced work, and the skepticism is well-placed. So it is worth being precise about what the model did and did not do here.

What it did not do: it did not decide what to measure, did not write the scales, did not choose the bins, and did not produce a single number that goes into the story without a human-readable quote attached.

What it did: read 4,094 documents and apply a written codebook to each one — the mechanical part of content analysis. That work is normally done by undergraduate coders working from exactly this kind of codebook, with exactly this kind of reliability check. The method here is standard content analysis; only the coder changed.

Why that is checkable in a way “the AI said so” is not:

- Every score is anchored to a verbatim quote, machine-verified against the source. There is no step that rests on the model’s summary of a document.
- The quotes make the classifier’s errors *visible*. If it mis-scores, the quote it cited is right there and obviously does not support the bin.
- Agreement is measured, not assumed (§6.2), and against humans, not just against itself (§6.3).
- The codebook is versioned, so a published number is traceable to the exact instructions that produced it.

The honest risk is not random error — the per-row check handles that. It is *correlated* error: a systematic misreading of one common phrasing, applied the same way to hundreds of documents. The v2 R1 near-miss described in §5 is exactly this failure mode. Three things guard against it: the sealed human set (§6.3), which would show a systematic gap; the requirement that every batch report its hard calls, which is how R1, R2, R10 and R12 were all caught; and the version comparison, which puts a size on how much a rule change moves the numbers.

What we recommend before publication: complete the human validation set (§6.3) and report the agreement figure in the story or a methods box, whatever it turns out to be. If the desk codes 200 syllabi by hand and agrees with the machine on the great majority, the finding stands on evidence rather than on anyone’s confidence in the tool. If it does not, we would rather find out ourselves.

9. Results

All 4,089 classifiable documents, covering 6,963 sections and 2,473 instructors. Run `scripts/analyze_policies.py` to regenerate every figure below; it prints each one by document, by section and by instructor, so the denominator is never implicit. Percentages here are by document unless stated.

Does the syllabus address it at all?

	documents	share	sections	share
Phones	1,606	39.7%	2,945	42.3%
Laptops	2,067	51.1%	3,564	51.2%
Generative AI	3,369	85.3%	5,800	83.3%

Most syllabi say nothing about phones at all. Three in five are silent. That is itself a finding, and it is the necessary context for everything below: the strictness figures describe the minority of instructors who chose to write a rule.

Among those that do address it, how strict?

	strict (4-5)	conditional (3)	permissive (1-2)
Phones	80.8%	9.2%	10.0%
Laptops	57.8%	3.7%	38.5%
Generative AI	50.9%	34.6%	14.5%

The headline

When a UT instructor writes a phone policy, four times in five it restricts or bans the phone outright. 546 syllabi (13.5% of all documents) prohibit phones with no stated exception; another 751 restrict them to permitted uses. Only 161 syllabi that mention phones are permissive about them.

Laptops are a different story, not a milder version of the same one. 1,092 syllabi — 26.9% — *require* students to have a laptop, and 38.5% of those that address laptops are permissive. Instructors are not moving against screens in general; they are moving against phones specifically, while increasingly depending on laptops.

AI is the most-addressed and least-settled. 85.3% of syllabi take a position, but they scatter: a quarter prohibit it outright, a third set conditions, and 2.8% could not be scored at all because the instructor left an unedited template with more than one option still in it — the highest unscorable rate of the three scales, and a sign of how recently this language was bolted on.

Figures that carry a caveat

- **AI strictness: 50.9%, band 49.1-56.7%.** 257 documents (7.6% of those stating an AI policy) sit on an unsettled rule boundary. See §7.7.
- **phone_required (5.3%) and laptop_required (26.9%) are overstated** by syllabi that require “a device” generically. See §7.3 and issue 6.
- **56 syllabi that ban smart glasses score 0 on both device scales**, so the strict end is understated by roughly that much. See §7.6.
- **Department-level breakouts must be computed per near-duplicate cluster**, not per document. See §3.

10. How to audit this yourself

Everything below runs against the database and needs no special access beyond the connection string.

Check a single score. Every row in `document_policies` has a document ID. Join to `documents` for `source_url`, open the syllabus, and read the sentence in `phone_evidence`, `laptop_evidence`, or `ai_evidence`. If the quote is not there or does not support the score, that is a finding — please report it.

Check a random sample. Pull 20 rows at random, verify each. Twenty rows is about ten minutes and is enough to notice a systematic problem.

Re-run the quote verification from scratch. `scripts/load_classifications.py` performs the check described in §6.1. It can be run against the raw classifier output at any time and will report every mismatch.

Re-code a batch by hand. `prompts/policy-classification-v7.md` is the full codebook. Anyone can apply it to 25 syllabi and compare.

Re-run everything from the source. `scripts/build_manifest.py`, `fetch_pdfs.py`, `fetch_simple_syllabus.py`, `extract_text.py` and `load_database.py` rebuild the corpus from UT's website. The syllabus site is public and unauthenticated.

Reported figures must name their codebook version (`prompt_version`) and their denominator (§3). A percentage without both is not reproducible.